References*

David Meyer

dmm@1-4-5.net

Last update: October 26, 2018

1 Introduction

This document contains references for the machine learning and networking parts of Project Number HE201707000. These references, in non-alphabetical order, are: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51]. Note that these references do not in general contain introductory references such as [52].

*This document is the final deliverable for Project Number HE2017070001.

References

[1] R. Istrate, A. Cristiano Innocenza Malossi, C. Bekas, and D. Nikolopoulos, "Incremental Training of Deep Convolutional Neural Networks," *ArXiv e-prints*, Mar. 2018.

[2] A. Valadarsky, M. Schapira, D. Shahaf, and A. Tamar, "A Machine Learning Approach to Routing," *ArXiv e-prints*, Aug. 2017.

[3] R. K. Ahuja, K. Mehlhorn, J. Orlin, and R. E. Tarjan, "Faster algorithms for the shortest path problem," *J. ACM*, vol. 37, pp. 213–223, Apr. 1990.

[4] O. San, R. Maulik, and M. Ahmed, "An artificial neural network framework for reduced order modeling of transient flows," *ArXiv e-prints*, Feb. 2018.

[5] F. J. Pineda, "Generalization of back-propagation to recurrent neural networks," *Phys. Rev. Lett.*, vol. 59, pp. 2229–2232, Nov 1987.

[6] F. Geyer and G. Carle, "Learning and generating distributed routing protocols using graph-based deep learning," in *Proceedings of the 2018 Workshop on Big Data Analytics and Machine Learning for Data Communication Networks, Big-DAMA@SIGCOMM 2018, Budapest, Hungary, August 20, 2018*, pp. 40–45, 2018.

[7] F. Scarselli, M. Gori, A. C. Tsoi, M. Hagenbuchner, and G. Monfardini, "The graph neural network model," *Trans. Neur. Netw.*, vol. 20, pp. 61–80, Jan. 2009.

[8] R. Dey and F. M. Salem, "Gate-Variants of Gated Recurrent Unit (GRU) Neural Networks," *ArXiv e-prints*, Jan. 2017.

[9] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, "Dropout: A simple way to prevent neural networks from overfitting," *J. Mach. Learn. Res.*, vol. 15, pp. 1929–1958, Jan. 2014.

[10] S. Andradóttir, D. P. Heyman, and T. J. Ott, "Variance reduction through smoothing and control variates for markov chain simulations," *ACM Trans. Model. Comput. Simul.*, vol. 3, pp. 167–189, July 1993.

[11] W. Genders and S. Razavi, "Using a Deep Reinforcement Learning Agent for Traffic Signal Control," *ArXiv e-prints*, Nov. 2016.

[12] E. Walraven, M. T. Spaan, and B. Bakker, "Traffic flow optimization," *Eng. Appl. Artif. Intell.*, vol. 52, pp. 203–212, June 2016.

[13] O. Nachum, S. Gu, H. Lee, and S. Levine, "Near-Optimal Representation Learning for Hierarchical Reinforcement Learning," *ArXiv e-prints*, Oct. 2018.

[14] M. Wang, Y. Cui, X. Wang, S. Xiao, and J. Jiang, "Machine Learning for Networking: Workflow, Advances and Opportunities," *ArXiv e-prints*, Sept. 2017.

[15] G. Stampa, M. Arias, D. Sanchez-Charles, V. Muntes-Mulero, and A. Cabellos, "A deep-reinforcement learning approach for software-defined networking routing optimization," 09 2017.

[16] N. Tao, N. Tao, J. Baxter, and L. Weaver, "A multi-agent, policy-gradient approach to network routing," *IN: PROC. OF THE 18TH INT. CONF. ON MACHINE LEARNING*, pp. 553–560, 2001.

[17] T. P. Lillicrap, J. J. Hunt, A. Pritzel, N. Heess, T. Erez, Y. Tassa, D. Silver, and D. Wierstra, "Continuous control with deep reinforcement learning," 09 2015.

[18] K. Sohn, H. Lee, and X. Yan, "Learning structured output representation using deep conditional generative models," in *Advances in Neural Information Processing Systems 28* (C. Cortes, N. D. Lawrence, D. D. Lee, M. Sugiyama, and R. Garnett, eds.), pp. 3483–3491, Curran Associates, Inc., 2015.

[19] J. Peters, J. M. Mooij, D. Janzing, and B. Schölkopf, "Identifiability of causal graphs using functional models," in *Proceedings of the Twenty-Seventh Conference on Uncertainty in Artificial Intelligence*, UAI'11, (Arlington, Virginia, United States), pp. 589–598, AUAI Press, 2011.

[20] G. Tucker, S. Bhupatiraju, S. Gu, R. E. Turner, Z. Ghahramani, and S. Levine, "The Mirage of Action-Dependent Baselines in Reinforcement Learning," *ArXiv e-prints*, Feb. 2018.

[21] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, "High-Dimensional Continuous Control Using Generalized Advantage Estimation," *ArXiv e-prints*, June 2015.

[22] C. Wu, A. Rajeswaran, Y. Duan, V. Kumar, A. M. Bayen, S. Kakade, I. Mordatch, and P. Abbeel, "Variance Reduction for Policy Gradient with Action-Dependent Factorized Baselines," *ArXiv e-prints*, Mar. 2018.

[23] S. Gu, T. Lillicrap, Z. Ghahramani, R. E. Turner, and S. Levine, "Q-Prop: Sample-Efficient Policy Gradient with An Off-Policy Critic," *ArXiv e-prints*, Nov. 2016.

[24] V. M. Aleksandrov, V. I. Sysoyev, and V. V. Shemeneva, "Stochastic optimization," *Engineering Cybernetics*, vol. 5, pp. 11–16, 1968.

[25] J. Baxter and P. L. Bartlett, "Infinite-horizon policy-gradient estimation," *J. Artif. Int. Res.*, vol. 15, pp. 319–350, Nov. 2001.

[26] D. Pfau and O. Vinyals, "Connecting Generative Adversarial Networks and Actor-Critic Methods," *ArXiv e-prints*, Oct. 2016.

[27] V. Mnih, A. Puigdomènech Badia, M. Mirza, A. Graves, T. P. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, "Asynchronous Methods for Deep Reinforcement Learning," *ArXiv e-prints*, Feb. 2016.

[28] A. Stooke and P. Abbeel, "Accelerated Methods for Deep Reinforcement Learning," *ArXiv e-prints*, Mar. 2018.

[29] V. R. Konda and J. N. Tsitsiklis, "Actor-critic algorithms," in *Advances in Neural Information Processing Systems 12* (S. A. Solla, T. K. Leen, and K. Müller, eds.), pp. 1008–1014, MIT Press, 2000.

[30] E. Greensmith, P. L. Bartlett, and J. Baxter, "Variance reduction techniques for gradient estimates in reinforcement learning," *J. Mach. Learn. Res.*, vol. 5, pp. 1471–1530, Dec. 2004.

[31] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal Policy Optimization Algorithms," *ArXiv e-prints*, July 2017.

[32] R. J. Williams, "Simple statistical gradient-following algorithms for connectionist reinforcement learning," *Machine Learning*, vol. 8, no. 3-4, pp. 229–256, 1992.

[33] J. Peters and S. Schaal, "Policy gradient methods for robotics," in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, (Beijing, China), 2006.

[34] D. Silver, G. Lever, N. Heess, T. Degris, D. Wierstra, and M. Riedmiller, "Deterministic policy gradient algorithms," in *Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32*, ICML'14, pp. I–387–I–395, JMLR.org, 2014.

[35] S. Kakade and J. Langford, "Approximately optimal approximate reinforcement learning," in *Proceedings of the Nineteenth International Conference on Machine Learning (ICML 2002)* (C. Sammut and A. Hoffman, eds.), (San Francisco, CA, USA), pp. 267–274, Morgan Kauffman, 2002.

[36] J. Schulman, S. Levine, P. Moritz, M. I. Jordan, and P. Abbeel, "Trust region policy optimization," *CoRR*, vol. abs/1502.05477, 2015.

[37] J. Schulman, S. Levine, P. Moritz, M. I. Jordan, and P. Abbeel, "Trust region policy optimization," 02 2015.

[38] J. Schulman, N. Heess, T. Weber, and P. Abbeel, "Gradient estimation using stochastic computation graphs," 06 2015.

[39] I. Goodfellow, "NIPS 2016 tutorial: Generative adversarial networks," 01 2017.

[40] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," 06 2017.

[41] D. Bahdanau, K. Cho, and Y. Bengio, "Neural machine translation by jointly learning to align and translate," 09 2014.

[42] A. Graves, "Generating sequences with recurrent neural networks," 08 2013.

[43] G. Shen, L. Li, Y. Zhang, W. Chen, S. K. Bose, and M. Zukerman, "Machine Learning-Assisted Least Loaded Routing to Improve Performance of Circuit-Switched Networks," *ArXiv e-prints*, Apr. 2018.

[44] G. Stampa, M. Arias, D. Sanchez-Charles, V. Muntes-Mulero, and A. Cabellos, "A Deep-Reinforcement Learning Approach for Software-Defined Networking Routing Optimization," *ArXiv e-prints*, Sept. 2017.

[45] M. Nazari, A. Oroojlooy, L. V. Snyder, and M. Takáč, "Reinforcement Learning for Solving the Vehicle Routing Problem," *ArXiv e-prints*, Feb. 2018.

[46] M. Kavalerov, Y. Likhacheva, and Y. Shilova, "A reinforcement learning approach to network routing based on adaptive learning rates and route memory," in *SoutheastCon 2017*, pp. 1–6, March 2017.

[47] S. Ayoubi, N. Limam, M. A. Salahuddin, N. Shahriar, R. Boutaba, F. E. Solano, and O. M. C. Rendon, "Machine learning for cognitive network management," *IEEE Communications Magazine*, vol. 56, no. 1, pp. 158–165, 2018.

[48] R. Boutaba, M. A. Salahuddin, N. Limam, S. Ayoubi, N. Shahriar, F. E. Solano, and O. M. C. Rendon, "A comprehensive survey on machine learning for networking: evolution, applications and research opportunities," *J. Internet Services and Applications*, vol. 9, no. 1, pp. 16:1–16:99, 2018.

[49] M. Polese, R. Jana, V. Kounev, K. Zhang, S. Deb, and M. Zorzi, "Machine Learning at the Edge: A Data-Driven Architecture with Applications to 5G Cellular Networks," *ArXiv e-prints*, Aug. 2018.

[50] F. Musumeci, C. Rottondi, A. Nag, I. Macaluso, D. Zibar, M. Ruffini, and M. Tornatore, "A Survey on Application of Machine Learning Techniques in Optical Networks," *ArXiv e-prints*, Mar. 2018.

[51] C. Streiffer, H. Chen, T. Benson, and A. Kadav, "DeepConfig: Automating Data Center Network Topologies Management with Machine Learning," *ArXiv e-prints*, Dec. 2017.

[52] R. S. Sutton and A. G. Barto, *Introduction to Reinforcement Learning.* Cambridge, MA, USA: MIT Press, 1st ed., 1998.