

LISP-CONS

A Mapping Database Service

David Meyer, Dino Farinacci, Vince Fuller, Darrel
Lewis, Scott Brim, Noel Chiappa

NANOG 41

October, 2007

<http://www.1-4-5.net/~dmm/talks/NANOG41/cons>

Agenda

- Brief Intro
- Design Considerations
- Brief Definitions
- How CONS Works
- What We've Learned
- Questions/Comments?

What is LISP?


- Locator/ID Separation Protocol (LISP)
 - draft-farinacci-lisp-03.txt
- Creates two namespaces: **IDs** and **Locators**
- Why do this?
 - Improve site multihoming
 - Improve ISP Traffic Engineering
 - Reduce site renumbering costs
 - Reduce size of core routing tables
 - PI for all?
 - Some form of mobility?

Locator/ID Split?

- The idea here is that the IP address is overloaded
 - It encodes both location in the topology (locator) and the identity of the user of the address
- The **locator** role is used by the routing system
- The **identity** role is used by upper layer protocols
 - e.g., TCP pseudo-header
- **Problem:** Since we want locators to aggregate topologically, and since identifiers are usually allocated on organizational boundaries, it is difficult (impossible?) to get one number space to efficiently serve both purposes.
 - There are other issues as well, including
 - The expected lifetime of a name (don't want to reconfigure...)
 - Who has control over the name(s)?
 - ...

Locator/ID Split?

- One solution: split the functions -- This is at the heart of the Locator/ID split idea
 - So how might we achieve this?
- Architecturally, we might try to “Jack-up” the existing IP layer



Implementing a Locator/ID Split

- There are two main ways to engineer a Loc/ID split
- **Rewriting**
 - If you have enough address space (e.g., IPv6), you could use the lower 64 bits as an identifier, and the upper 64 bits as a locator, and rewrite the locator at the border
 - This is the basis of O'Dell's 8+8/GSE scheme
 - Credit to Bob Smart and Dave Clark on this one too
- **Map-n-Encap**
 - You could also put another header on the packet, and make the inner header carry the IDs and the outer header carry the locators
 - LISP is an instance of this approach
 - Credit to Bob Hinden & Steve Deering on map-n-encap...

Loc/ID Split in practice

IPv6: 2001:0102:0304:0506:1111:2222:3333:4444

Locator ID

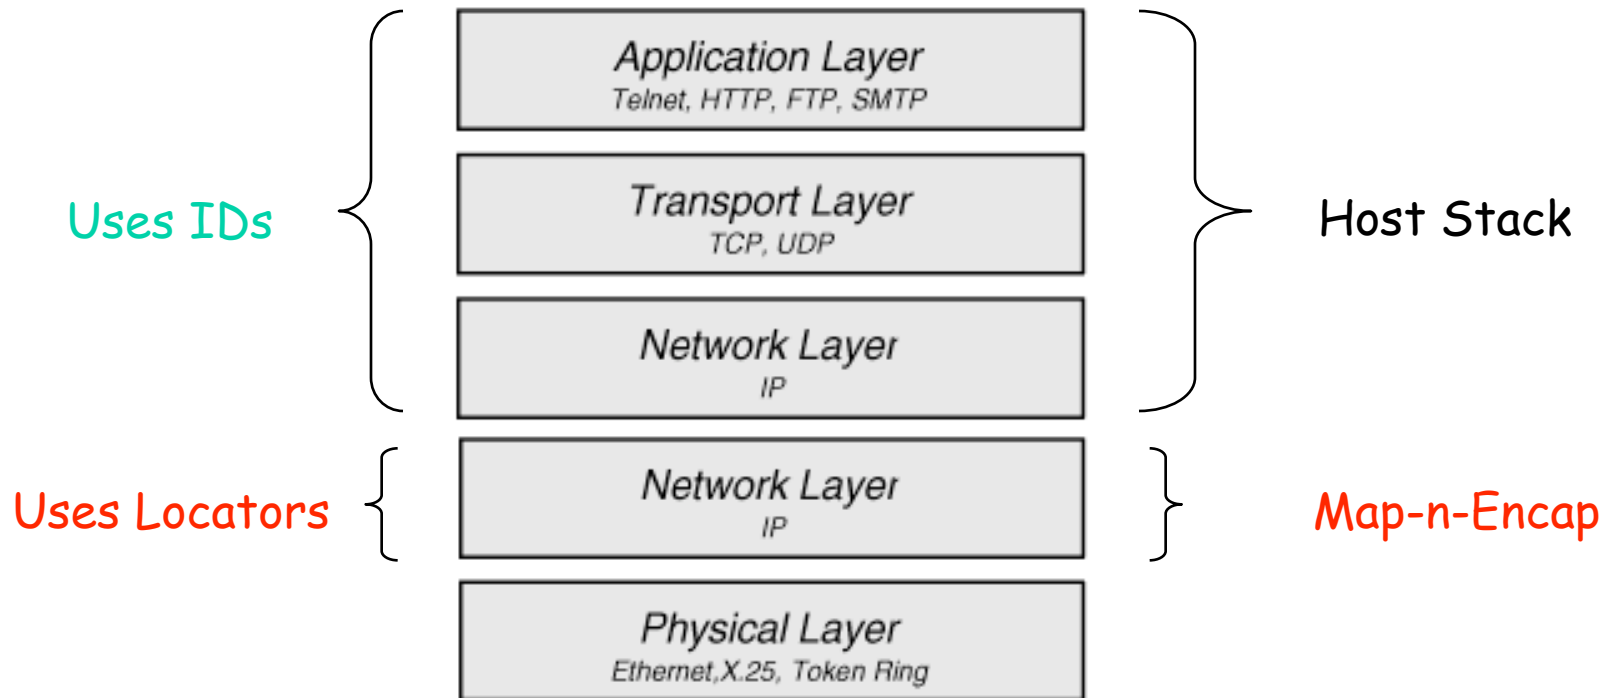
IPv4

209.131.36.158.10.0.0.1

:

Locator ID

LISP is a Jack-Up



LISP Parts

- Data-plane
 - Design for encapsulation and tunnel router placement
 - Design for locator reachability
 - Data triggered mapping service
- Control-plane
 - Design for a scalable mapping service
 - This talk is about LISP Control-planes

LISP Variants

- LISP 1
 - Rutable IDs over existing topology to probe for mapping reply
- LISP 1.5
 - Rutable IDs over another topology to probe for mapping reply
- LISP 2
 - EIDs are not routable and mappings are in DNS
- LISP 3
 - EIDs are not routable, mappings obtained using new mechanisms (DHTs perhaps, LISP-CONS, NERD, APT)

Data-Plane Mapping
Control-Plane Mapping

Quick LISP Terms

- Endpoint Identifiers (**EIDs**)
 - IDs for host-use and only routeable in source and dest sites
 - Can be out of PA or PI address space
- Routing **Locators** (**RLOCs**)
 - Routable addresses out of PA address space
- Ingress Tunnel Router (ITR)
 - Device in source-site that prepends LISP header with RLOCs
- Egress Tunnel Router (ETR)

LISP Control-Plane

- Build a large distributed mapping database service
- Scalability paramount to solution
- How to scale: $(\text{state} * \text{rate})$
- If both factors large, we have a problem
 - `state` will be "large" ($O(10^{10})$ hosts)
 - Aggregate EIDs into EID-prefixes to reduce state
 - So `rate` must be small
 - Make mappings have "subscription time" frequency
 - i.e., we expect such mappings to change with low frequency
 - And **no reachability information in the mapping database**

Some Questions for a LISP Control-Plane

- Where to put the mappings?
- How to find the mappings?
- Is it a push model?
- Is it a pull model?
- Do you use secondary storage?
- Do you use a cache?
- What about securing the mapping entries?
- What about protecting infrastructure from DOS-attacks?
- What about controlling packet loss and latency?

LISP Control-Plane

"Push doesn't scale, caching doesn't
scale, pick one"

LISP-CONS

- LISP-CONS is a hybrid approach
- Push EID-prefixes (**but not mappings**) at upper levels of hierarchy
- Pull from lower levels of hierarchy
- Mappings stored at lower-levels
 - Requests get to where the mappings are
 - Replies are returned
 - This is a crucial point as we'll see in a bit
- Getting to the lower-levels via pushing of EID-prefixes
- LISP-CONS is a mapping system for LISP 3.0

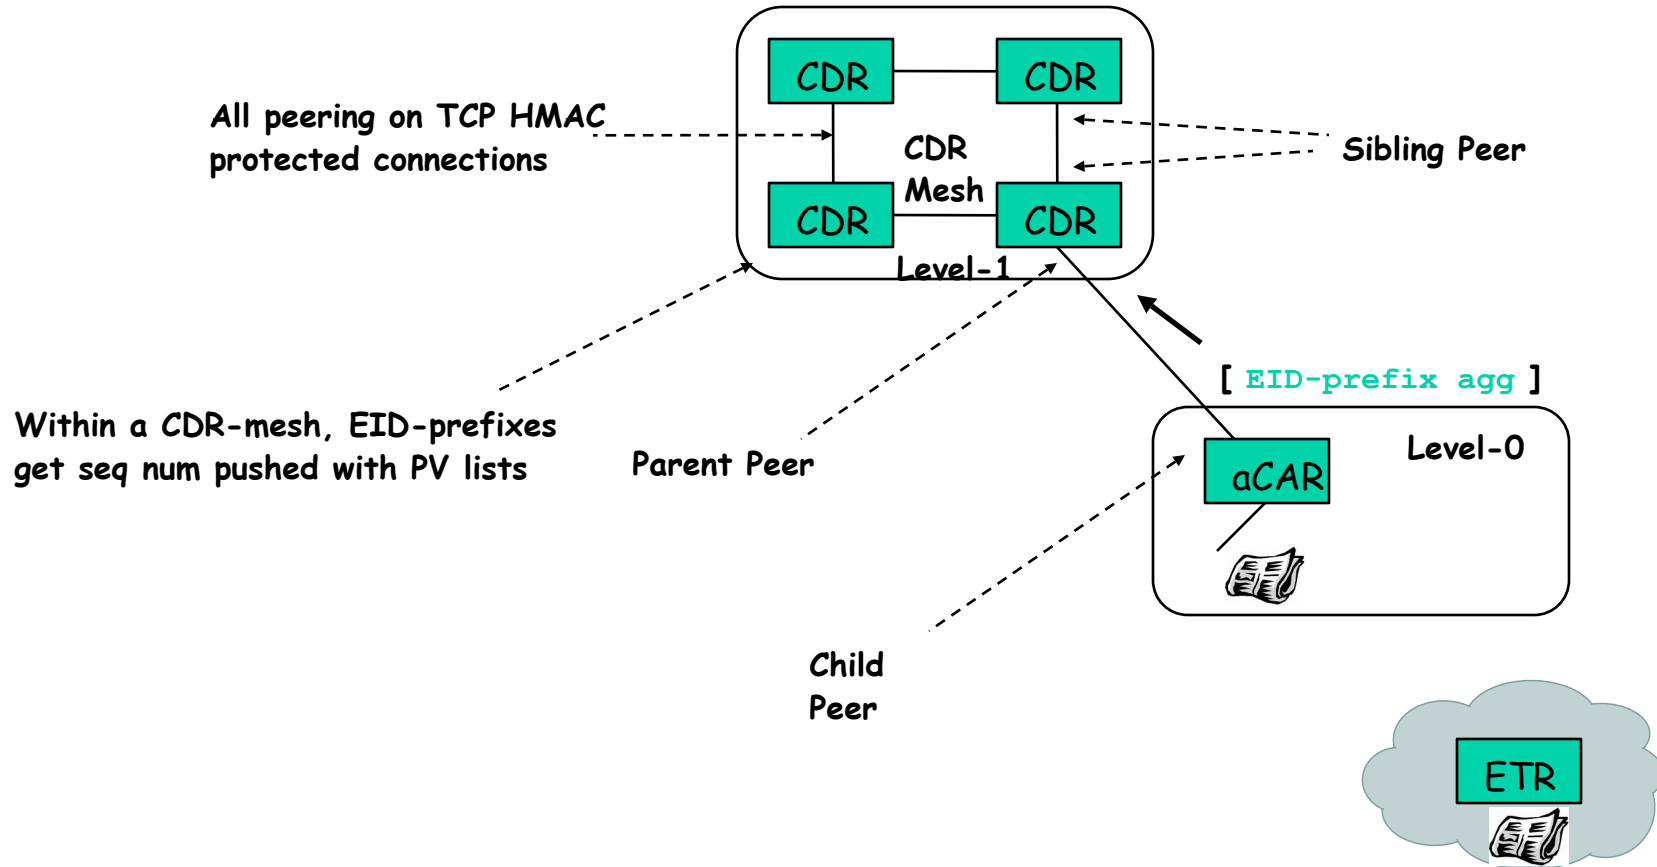
LISP-CONS

- We can get good EID-prefix aggregation
 - If hierarchy based on EID-prefix allocation and not topology
 - Then build a logical topology based on the EID-prefix allocation
- Map-Requests routed through logical hierarchy
 - Key is the EID
- Map-Reply returned to originator
 - With mapping record {EID-prefix, RLOC-set}

LISP-CONS Network Elements

- Content Access Routers (CARs)
 - Querying-CARs
 - Generate Map-Requests on behalf of ITRs
 - Answering-CARs
 - Hold authoritative mappings at level-0 of hierarchy
 - Aggregate only EID-prefix upwards
 - Respond with Map-Replies
- Content Distribution Routers (CDRs)
 - Push around EID-prefixes with level-1 to n of hierarchy
 - Aggregate EID-prefix upwards
 - Advertise EID-prefixes in a mesh topology within level
 - Forward Map-Requests and Map-Replies

LISP-CONS -- Peering



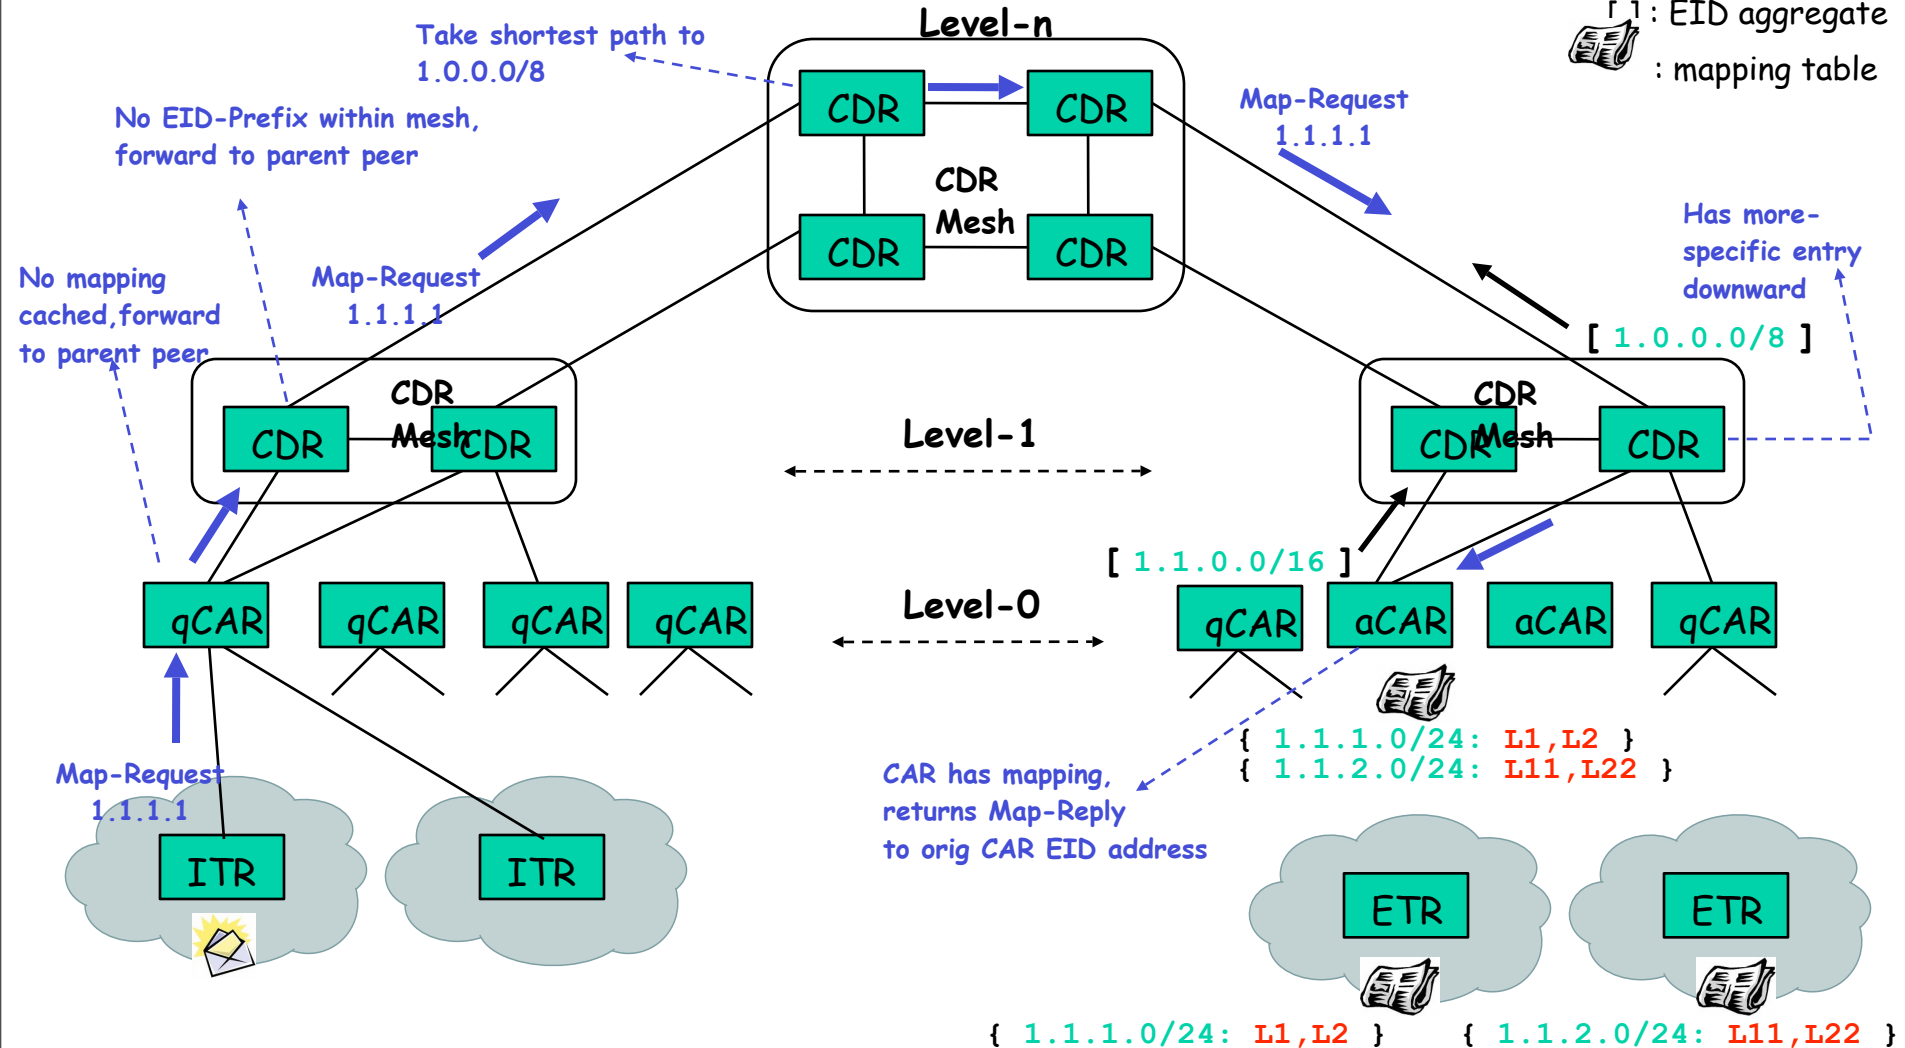
Here's how it works

Legend:

{ } : mapping entry

[] : EID aggregate

 : mapping table



What We've Learned

- We wanted to optimize aggregatability of EID prefixes
 - That led to the design in which only EID prefixes were pushed around at the higher levels (but not the mappings themselves)
 - We were concerned about the **rate*state** product
- However, some SPs articulated another dimension
 - **Latency**
 - So you have to tradeoff **rate, state, and latency**
 - If you push, you wind up with the whole database in network elements (**state**)
 - If you pull, you incur **latency**
 - If you try to do mobility, you get lots of updates (**rate**)

What We've Learned

- Current thinking is that a different hybrid approach might be most feasible
- Push the whole mapping table around in the "CDR" level
- ITRs pull mappings from the "CAR" level
- This has a few nice properties:
 - You can get the whole mapping table
 - If you happen to want it
 - Latency is reduced because you don't have to traverse the whole hierarchy to retrieve the mappings

Drafts

- LISP

- `draft-farinacci-lisp-03.txt`

- CONS

- `draft-meyer-lisp-cons-02.txt`

Questions/Comments?

Thanks!